

## Homework #6: Building a Distributed Map/Reduce Framework

Checkpoint due Sunday, April 26 at 11:59 p.m.

Due Thursday, April 30 at 11:59 p.m.

In this assignment you will implement a distributed map/reduce framework. Map/reduce is a programming model for processing large data sets, typically implemented by distributing a computation across multiple worker servers where the data is stored. The computation is provided as a plug-in to the framework, and is specified as a separate map task and a reduce task. A client submits the two tasks to a master server, which is in charge of distributing the tasks across the worker servers in the system and returning the result back to the client when it is complete.

Your goals in this assignment are to:

- Gain a deep understanding of map/reduce and the challenges of building a distributed system by implementing a map/reduce framework.
- Learn how to use a map/reduce framework by implementing a simple map/reduce task that uses your framework.
- Practice network programming using **Sockets**.
- Practice parallel and concurrent programming in Java.

For this assignment we encourage you to discuss map/reduce and high-level details of your solution with the course staff and with your classmates, but you may not share your code with other students and must submit your own solution for the assignment.

This assignment includes a small checkpoint deadline (Sunday, April 26 at 11:59 p.m.), and is due Thursday, April 30 at 11:59 p.m. The usual course late policy applies for the overall deadline, but you may **not** use any late days for the checkpoint deadline.

### Designing and implementing your map/reduce framework

---

Your framework must match the basic architecture described in lecture. Specifically, your framework should consist of three main components: a client, a master server, and multiple worker servers. The client simply submits the map and reduce tasks to the master server and waits for the master server to confirm that the computation is complete.

A map/reduce computation consists of two phases: a map phase and a reduce phase. In the map phase the master server distributes the map task across a set of worker servers.

The master server specifies on which subset of the data each map worker should execute the map task based on what data the worker has access to. The map workers then execute the map task on that subset of the data and write their *intermediate results* to a file on their local disk.

After all map workers have completed the map task the master server begins the reduce phase, distributing the reduce task across worker servers—sometimes but not necessarily the same set of workers that executed the map task. The master server specifies on which subset of the intermediate results each reduce worker should execute the reduce task. Each reduce worker begins by obtaining the specified subset of intermediate results from the map workers, a step known as *the shuffle*. In the shuffle it's often necessary for every reduce worker to obtain some intermediate results from every map worker; in practice implementing the shuffle efficiently is one of the most complex aspects of a map/reduce framework. At the end of the shuffle, each reduce worker has a disjoint subset of the intermediate results. Each reduce worker then executes the reduce task and saves the final key/value pair results in its file system, notifying the master server when it is complete. When the master server is notified that all reduce workers are complete, the master server notifies the client that the computation is complete and specifies the file locations that store the final results.

If a worker server crashes or otherwise fails to return a result, the master server distributes the map or reduce task to another worker. In a typical map/reduce framework the data is redundant in a distributed file system, and the map/reduce framework should produce the correct result any time the full data set is available on non-failed servers. In the case that the full set of data is unavailable, there is no other choice but to return the result based on the available data.

### A faux distributed storage system

To achieve efficient operation, map/reduce is typically coupled closely to an implementation of a distributed storage system, which stores both the source data and the results of the map/reduce computation. For map/reduce to be efficient the map/reduce worker servers are usually servers in the distributed storage system; each map worker usually just computes the map task for a subset of the data it locally stores for the distributed storage system, and each reduce worker stores the map/reduce results in the distributed storage system, but in files stored at the reduce worker's local file system.

In this assignment, however, we have not provided a distributed storage system and you should **not** implement your own. Instead, we have provided sample data partitioned much as it might be partitioned within a distributed storage system, into separate file sets that could each be replicated and served from multiple storage servers. For testing purposes, you may copy all of this data to any server you use to test your solution, but your worker

server should mimic the behavior of a server within the storage system and only access the subset of the data specified as available to that server.

For this assignment our expectation is that each partition is locally available at one or more worker servers. You may assume that the worker-to-partition association is static and known by all workers and the master server when those programs start. When your master server distributes a map task among worker servers, it should assign each map worker to execute the task on a subset of the data that worker is hosting.

To simplify your work we have provided a `Partition` class that, given a partition name (e.g., "7") and worker name, allows you to iterate over all files stored in that partition for that worker. This `Partition` class assumes the files are available in a directory layout specific to our faux distributed storage system, and we have provided sample partitioned data in your Eclipse project.

### **The MapTask, ReduceTask, and Emitter interfaces**

We have also provided a plug-in interface for computations for your map/reduce framework, as well as a sample word count computation you may use to test your implementation. Like the map/reduce example from class, our plug-in interface requires the results of both map and reduce tasks to be `String/String` key/value pairs.

Based on the example map/reduce computation from class, our `MapTask` and `ReduceTask` interfaces should be mostly self-explanatory. These interfaces use an `Emitter` interface that allows `MapTask` and `ReduceTask` implementations to communicate their results—each result being a single key/value pair—to the framework, which may then process those results as needed.

### **Starting this homework assignment**

The work described in this section is due by the checkpoint deadline, Sunday, April 26 at 11:59 p.m.

Start by completing the diagram on the last page of this document to describe the network interactions between servers in your map/reduce system on a simple word-count computation. We do not require precise notation, but use a notation similar to sequence diagrams to describe the communication between servers: draw an arrow from one server to another to indicate a network message, and label the arrow with the contents of the message. Your message descriptions do not need to be precise or formal. E.g., if a worker sends a message to the master server to indicate that the worker has completed some task A, you can just label the arrow "Done with A". Only indicate network communication between servers; you do not need to describe computation or interactions between internal server components.

For the example word-count computation suppose the worker servers are storing three

replicated partitions (P1, P2, and P3) each containing a single file:

- P1: Red Mean TwoTwo
- P2: Mine Mine Mean
- P3: Red TwoTwo Mean

For the purposes of this diagram, assume that applying some arbitrary hash function,  $h$ , the above words gives these values:

- $h(\text{"Red"}) = 0$
- $h(\text{"Mean"}) = 1$
- $h(\text{"TwoTwo"}) = 2$
- $h(\text{"Mine"}) = 3$

You may hand-write and scan or use a computer program to complete the diagram; submit your work as `mapreduce_diagram.pdf` as part of the checkpoint submission.

After you have completed the diagram, you should begin work on your map/reduce framework. By the time of the checkpoint deadline, you should ensure that your `WorkerServer` can receive and execute an arbitrary `MapTask`. To accomplish this goal you must implement the `WorkerServer#run()` method and implement any other components (such as implementing the `Emitter` interface) you need to execute a `MapTask`. To verify you have correctly completed this task, we also recommend that you partially complete the `MasterServer#run()` method so that you can send `MapTask` implementations to your `WorkerServer` implementation for execution.

### Completing your map/reduce framework

For the final portion of the framework, you must complete the `AbstractClient#execute()` method and the `MasterServer#run()` method, as well as design and implement any other components you need for `MapTask` and `ReduceTask` implementations to obtain data from and communicate results to the framework. Debugging this portion of the homework can be extremely tricky so make sure you:

- Log your exceptions, especially in any `Runnable#run()` and `Callable#call()` methods. This way you know where and how your code went wrong in different concurrent operations.

- Make use of the **synchronized** keyword to protect yourself from race conditions. A concurrent hash map does not protect against all race conditions.
- Store your intermediate results in a human-readable format. This will make it simpler you to debug your results and us to grade your work.

When you are done, you should be able to run the **MasterServer**, **WorkerServer**, and a client (such as the **WordCountClient**) as separate Java programs. The client should submit a map task and reduce task to the master server, which should manage the map/reduce computation as described above. The map/reduce computation should be executed on all files (excluding replicas) in the faux distributed storage system. All communication between framework components must use network sockets: client to/from master server; master to/from worker servers; and pairs of worker servers. To clarify, even though workers may be simulated as threads in a single Java program, your framework must use network sockets to transmit intermediate results between map and reduce workers in the shuffle.

The master server should assign tasks to workers as follows:

- Each worker can complete its task given its subset of the data. This means that each worker should only be able to read and write to its own set of data and nothing else.
- The map/reduce computation obtains the correct result any time the full data set is available. If a worker becomes unavailable during a computation, the master server should account for the failure and redistribute the task to other available workers.

Our **MasterServer**, **WorkerServer**, and client implementations contain **main** methods that start separate master, worker, and client programs. When run without any command-line arguments these programs obtain a sample master and worker configuration from two Java properties files, **master.properties** and **worker.properties**, initializing master, worker, and client classes with host names and ports that you can use to create network sockets for inter-process communication. By default, the **WorkerServer** simulates multiple workers in a configuration by running multiple threads in a single Java program.

Inside the **worker.properties** file we've defined a sample configuration of four worker servers, with each worker server storing a subset of the data in the faux distributed storage system and each partition being stored redundantly by at least two workers. In this configuration your map/reduce framework should be able to successfully complete a map/reduce computation if any single worker server fails during the computation.

Alternatively, the configuration for the **MasterServer**, **WorkerServer**, and client programs may be specified as command line arguments. You may specify a configuration on the command line to enable the worker servers to run as separate processes (rather than as separate threads within a single process) to facilitate testing with worker failures.

## Using your map/reduce framework

---

When you have completed your map/reduce implementation, you must use your map/reduce framework by executing our sample computation on your framework and then implement a new map/reduce computation of your own. The tasks are described below.

### A sample word count computation

We have provided a sample map/reduce computation to count all occurrences of all words in a corpus of data, as a `WordCountMapTask`, a `WordCountReduceTask` and a `WordCountClient` that reads our master and worker configuration files. This computation was also provided as an example in lecture, and you should use it to test your framework. For reference, the word “a” appears 9,976 times and “and” appears 16,299 times in the sample data set we provide in your Git repository.

### Suggesting words based on word prefixes

Once you’ve tested your solution using our `WordCountClient` example, write a map/reduce task to help determine the best word completions for word prefixes. Specifically, your map/reduce task should analyze a corpus of data and output a *prefix/word* key/value pair for each prefix that appears in the corpus, where the word associated with a given prefix is the word that most frequently completes that prefix in the corpus.

For example, consider the prefix “a”. Many words in a given corpus of data will have the prefix “a”, since many words begin with the letter “a”. For our corpus, the most common word beginning with “a” is the word “and”, so the output of your map/reduce task output should include the key/value pair `a/and`. Similarly, the prefix “heav” could start one of several words—heavy, heaven, heavenly, etc. In our corpus, “heaven” appears the most frequently (87 times to be exact) of any other word that begins with the prefix “heav”, so your map/reduce task output should include the key/value pair `heav/heaven`.

## Evaluation

---

Your solution should meet the following minimal requirements:

- By the checkpoint deadline you must submit solutions for the tasks described in the section “Starting this homework assignment.”
- Map workers should write their (human-readable) intermediate results directly to the file system. Reduce workers should write the final (human-readable) key/value pairs to the file system (in the `final_results` directory) and return the locations of

where the final key/value pair results are stored. Your reduce workers must generate unique file names for each computation so that the framework does not overwrite the results of previous computations; you can generate unique file names by including the time stamp of the computation as part of the file name. The master server should aggregate the results locations from each reduce worker and return this information to the client, marking the end of the map/reduce computation. It would be wise to have your client print these results to console so you can easily determine if your results are correct.

- You must use network sockets to implement all communication between the client and master, between the master and each worker, and between pairs of worker servers. During the shuffle, workers must communicate with other workers using network sockets (when we test your implementation, we will run each worker on a separate computer, so your implementation won't work otherwise).
- Your framework should be designed to work with an arbitrary (but fixed) number of worker servers. For example, if we add a worker server to our `workers.properties` configuration file, your system should correctly implement map/reduce including that worker server, without requiring any programmatic changes to your solution.
- Your framework should be robust to simple worker failures, and should compute the correct result as long as the full set of data is available on non-failed worker servers. Your framework is allowed to compute an incorrect result for a map/reduce computation if enough workers fail to make some partition data unavailable.
- Follow standard Java style guidelines and document your code with Javadocs where appropriate. FindBugs should report no bugs at the standard level (15).

Overall, this homework is worth 100 points. You will be graded as follows:

- Checkpoint completion: 10 points.
- Working map/reduce framework using separate master/worker programs and network communication when deployed on multiple servers: 50 points.
- Robustness of map/reduce framework when some worker servers are unavailable but the full data set is available: 10 points.
- Correct implementation of the word-prefix computation: 20 points.
- Javadocs, style, and FindBugs: 10 points.

Master	Worker 1	Worker 2	Worker 3	Worker 4	Description
	P1, P2	P1, P2	P3	P2, P3	Initial Contents
					<b>Start Map</b>  Master tells worker servers to start the map phase, providing the map function.
	mapResult <div></div>	mapResult <div></div>	mapResult <div></div>	mapResult <div></div>	<b>Workers Perform Map</b>  Workers execute the map task and store their results in the mapResult file. Then, they inform the master server that the map task has been completed.  If a worker server is not used, leave the box blank.
					<b>Start Shuffle</b>  Master tells worker servers to start the shuffle phase.
	shuffleResult <div></div>	shuffleResult <div></div>	shuffleResult <div></div>	shuffleResult <div></div>	<b>Workers Perform Shuffle</b>  Workers send the map results to each other, based on the hashcode of the key. Workers store these results in another file. Then, they inform the master server when they are done.  If a worker server is not used, leave the box blank.
					<b>Start Reduce</b>  Master tells worker servers to start the reduce phase, providing the reduce function.
	finalResult <div></div>	finalResult <div></div>	finalResult <div></div>	finalResult <div></div>	<b>Workers Perform Reduce</b>  Workers execute the reduce task, creating a new file with the final results of the map-reduce task.  Then, they inform the master server they are done, providing the file location of the result.  In this case, the file location is always the same for simplicity.